

ACSE-4 Project 2: A Smoothed Particle Hydrodynamics Solver

Introduction to SPH

Smoothed Particle Hydrodynamics (SPH) is a meshless method for solving the Navier-Stokes equation. As it is meshless and Lagrangian (reference frame follows the fluid motion), it is ideal for solving problems fluid flow with interfaces and free surfaces. In this task you are going to write an SPH-based simulator and use it to solve a wave generation problem. The equations listed below include some which are purely there for background context and some which are required within the SPH simulation code. Those which you will need to implement are surrounded by a box.

The basic idea of SPH is that you model individual points which follow the fluid and are each associated with a given mass of the fluid, m_i . In order to do differentiation, the discrete values of any field (e.g., pressure) associated with the points are converted into a continuous function using a smoothing kernel (hence the smoothed in the name):

$$\langle A \rangle(\mathbf{x}) = \sum_{j=1}^N m_j \frac{A_j}{\rho_j} W(\mathbf{x} - \mathbf{x}_j, h)$$

Where $\langle A \rangle(\mathbf{x})$ is the smoothed value of A at location \mathbf{x} , which is a vector, A_j is the value of A associated with particle j , ρ_j is the density of particle j and $W(\mathbf{x} - \mathbf{x}_j, h)$ is a smoothing kernel which depends upon the vector $\mathbf{x} - \mathbf{x}_j$ and a characteristic smoothing length h .

In SPH we are typically only interested in the values at the location of the points. We will thus use the following notation:

$$\langle A \rangle(\mathbf{x}_i) = \langle A \rangle_i = \sum_{j=1}^N m_j \frac{A_j}{\rho_j} W(\mathbf{r}_{ij}, h)$$

Where $\mathbf{r}_{ij} = \mathbf{x}_i - \mathbf{x}_j$. Also note that as this is a smoothing rather than an interpolation, in general $\langle A \rangle_i \neq A_i$. At any point there is therefore a choice that can be made as to whether to use the smoothed or nominal value at that point.

The smoothing kernel must possess the property that it integrates to 1 when integrated over its area (or volume in 3D). The form of this smoothing kernel is important. It should have compact support (i.e. it should go to zero at a finite distance) so that only the neighbouring particles are involved in the calculations. It should also be smooth and, ideally, close to normal in shape. A number of different kernels have been proposed, but in this project I wish you to use the cubic spline:

$$W(\mathbf{r}, h) = \frac{10}{7\pi h^2} \begin{cases} 1 - \frac{3}{2}q^2 + \frac{3}{4}q^3 & \text{if } 0 \leq q \leq 1 \\ \frac{1}{4}(2 - q)^3 & \text{if } 1 \leq q \leq 2 \\ 0 & \text{if } q > 2 \end{cases}$$

Where $q = \frac{|\mathbf{r}|}{h}$

Note that this kernel is symmetric and thus only depends on the distance between the points being considered. The prefactor in the equation will depend upon the number of dimensions and this equation as written is for 2D problems only.

If we want to use SPH to approximate the Navier-Stokes equation we need to have approximations for the partial derivatives of the variables of interest:

$$\nabla A_i \approx \sum_{j=1}^N m_j \frac{A_j}{\rho_j} \nabla W(\mathbf{r}_{ij}, h)$$

This is useful as $\nabla W(\mathbf{r}_{ij}, h)$ is known analytically, while all the other values are also known for each point. This means that calculating the gradient of a quantity is no more computationally expensive than calculating the average value at a point. For symmetric kernels such as we will be using in this project, we can further simplify this equation by using the following relationship:

$$\nabla W(\mathbf{r}_{ij}, h) = \frac{dW}{dr}(|\mathbf{r}_{ij}|, h) \mathbf{e}_{ij}$$

Where $\mathbf{e}_{ij} = \frac{\mathbf{r}_{ij}}{|\mathbf{r}_{ij}|}$ and $\frac{dW}{dr}(r, h)$ is a known function of r and h , which you will need to derive by differentiating the kernel (note that you must do this analytically, not numerically). Substituting this back into the previous equation gives:

$$\nabla A_i \approx \sum_{j=1}^N m_j \frac{A_j}{\rho_j} \frac{dW}{dr}(|\mathbf{r}_{ij}|, h) \mathbf{e}_{ij}$$

The problem with this equation is that it is neither symmetric nor anti-symmetric (the contribution of particle i on particle j is not the same as the contribution of particle j on particle i). (Anti)symmetry is desirable if we wish the system to explicitly conserve momentum and mass.

To achieve symmetric or anti-symmetric differentials we can use one of the following two relationships:

$$\rho \nabla A = \nabla(\rho A) - A \nabla \rho$$

$$\frac{\nabla A}{\rho} = \nabla \left(\frac{A}{\rho} \right) + \frac{A}{\rho^2} \nabla \rho$$

Substituting the previous relationship into these equations results in the following approximations:

$$\rho \nabla A \approx \sum_{j=1}^N m_j (A_j - A_i) \frac{dW}{dr}(|\mathbf{r}_{ij}|, h) \mathbf{e}_{ij}$$

$$\frac{\nabla A}{\rho} \approx \sum_{j=1}^N m_j \left(\frac{A_i}{\rho_i^2} + \frac{A_j}{\rho_j^2} \right) \frac{dW}{dr}(|\mathbf{r}_{ij}|, h) \mathbf{e}_{ij}$$

Approximating the Navier-Stokes Equation using SPH

We now have all the tools required to approximate the Navier-Stokes equation and the associated continuity equations:

$$\frac{D\mathbf{v}}{Dt} = -\frac{\nabla P}{\rho} + \frac{\mu}{\rho} \nabla^2 \mathbf{v} + \mathbf{g}$$

$$\frac{D\rho}{Dt} = -\rho \nabla \cdot \mathbf{v}$$

It should be noted that in a Lagrangian reference frame $\frac{D}{Dt} \equiv \frac{\partial}{\partial t}$ when the time derivatives are carried out within the reference frame of differential volumes of the moving fluid (which in the SPH context is equivalent to the reference frame of the particles).

We have thus far not dealt with approximations for second derivatives, which are required to approximate the shear stress term. We could do this by taking the 2nd derivative of the kernel, but this is not a good approximation unless there are a large number of particles in the support region of the kernel, which is computationally expensive. An alternative is to use the following relationship:

$$\nabla^2 \mathbf{v}_i = \nabla \cdot \nabla \mathbf{v}_i \approx \sum_{j=1}^N m_j \left(\frac{1}{\rho_i^2} + \frac{1}{\rho_j^2} \right) \frac{dW}{dr}(|\mathbf{r}_{ij}|, h) \nabla \mathbf{v}_{ij} \cdot \mathbf{e}_{ij}$$

We can make use of the fact that $\nabla \mathbf{v}_{ij} \cdot \mathbf{e}_{ij}$ is the velocity gradient in the ij direction. This can be approximated as the difference in the velocity in the ij direction divided by the distance over which this occurs. Combining this approximation with the previous gradient approximations results in the following approximations for the Navier-Stokes and continuity equations:

$$\begin{aligned} \frac{\partial \mathbf{v}_i}{\partial t} = \mathbf{a}_i &= - \sum_{j=1}^N m_j \left(\frac{P_i}{\rho_i^2} + \frac{P_j}{\rho_j^2} \right) \frac{dW}{dr}(|\mathbf{r}_{ij}|, h) \mathbf{e}_{ij} + \mu \sum_{j=1}^N m_j \left(\frac{1}{\rho_i^2} + \frac{1}{\rho_j^2} \right) \frac{dW}{dr}(|\mathbf{r}_{ij}|, h) \frac{\mathbf{v}_{ij}}{|\mathbf{r}_{ij}|} \\ &\quad + \mathbf{g} \\ \frac{\partial \rho_i}{\partial t} = D_i &= \sum_{j=1}^N m_j \frac{dW}{dr}(|\mathbf{r}_{ij}|, h) \mathbf{v}_{ij} \cdot \mathbf{e}_{ij} \end{aligned}$$

Note that the sums in the above equations MUST NOT include the combinations where $i = j$ as this will result in a division by zero!

These are the main two equations that you will need to solve. Note the dot product in the continuity equation. You will notice that you currently have 3 equations (momentum equations in the x and y directions and the continuity equation) and 4 unknowns (velocity in the x and y directions, pressure and density) at each point. There is therefore still a missing equation required to close the system.

The way we will be closing the system is to assume that it is slightly compressible and to use an equation of state, which relates the pressure to the density. In this work we will use the Tait equation, which is a very stiff equation of state:

$$P = B \left(\left(\frac{\rho}{\rho_0} \right)^\gamma - 1 \right)$$

Where $B = \frac{\rho_0 c_0^2}{\gamma}$. ρ_0 is the initial or reference density of the fluid and c_0 is a numerical speed of sound, which should NOT be the same as the actual speed of sound (the simulation would be prohibitively slow). The value of c_0 should be chosen to be about an order of magnitude larger than the highest sustained speed in the system (for your work you can use $c_0 = 20 \text{ m/s}$, though you should test its impact on the results). To make sure that the equation is stiff enough use $\gamma = 7$.

Time stepping in SPH

The above equations allow a particle's acceleration and rate of change of density to be calculated as a function of its neighbours positions, velocities and densities/pressures. We now need to be able to update the particle's position, velocity and density, with the updated density being used to calculate a new pressure. To save complexity we will be doing this explicitly. We could use a simple explicit Euler scheme:

$$\begin{aligned} \mathbf{x}_i^{t+1} &= \mathbf{x}_i^t + \Delta t \mathbf{v}_i^t \\ \mathbf{v}_i^{t+1} &= \mathbf{v}_i^t + \Delta t \mathbf{a}_i^t \\ \rho_i^{t+1} &= \rho_i^t + \Delta t D_i^t \end{aligned}$$

This, though, is only first order accurate and is likely to be not very stable unless very small time steps are used. A scheme which is second order accurate in time is a predictor-corrector scheme:

Half-step

$$\mathbf{x}_i^{t+1/2} = \mathbf{x}_i^t + 0.5\Delta t \mathbf{v}_i^t$$

$$\mathbf{v}_i^{t+\frac{1}{2}} = \mathbf{v}_i^t + 0.5\Delta t \mathbf{a}_i^t$$

$$\rho_i^{t+1/2} = \rho_i^t + 0.5\Delta t D_i^t$$

Full-step

$$\bar{\mathbf{x}}_i^{t+1/2} = \mathbf{x}_i^t + 0.5\Delta t \mathbf{v}_i^{t+1/2}$$

$$\bar{\mathbf{v}}_i^{t+\frac{1}{2}} = \mathbf{v}_i^t + 0.5\Delta t \mathbf{a}_i^{t+\frac{1}{2}}$$

$$\bar{\rho}_i^{t+1/2} = \rho_i^t + 0.5\Delta t D_i^{t+1/2}$$

$$\mathbf{x}_i^{t+1} = 2 \bar{\mathbf{x}}_i^{t+1/2} - \mathbf{x}_i^t$$

$$\mathbf{v}_i^{t+1} = 2 \bar{\mathbf{v}}_i^{t+1/2} - \mathbf{v}_i^t$$

$$\rho_i^{t+1} = 2 \bar{\rho}_i^{t+1/2} - \rho_i^t$$

The added complexity with the predictor-corrector scheme is that 2 values for each of the variables need to be stored (the initial value and the half-step value) and the neighbour searching needs to be based on the appropriate \mathbf{x} values.

We also need to choose a time step that will ensure stability. In this system the following time step constraints need to be satisfied:

$$\Delta t_{CFL} = \min \left\{ \frac{h}{|\mathbf{v}_{ij}|} \right\}$$

$$\Delta t_F = \min \left\{ \sqrt{\frac{h}{|\mathbf{a}_i|}} \right\}$$

$$\Delta t_A = \min \left\{ \frac{h}{c_0 \sqrt{(\rho/\rho_0)^{\gamma-1}}} \right\}$$

$$\Delta t = C_{CFL} \min\{\Delta t_{CFL}, \Delta t_F, \Delta t_A\}$$

With an appropriate value of C_{CFL} being in the range 0.1 - 0.3 (top end of the range is quicker, but less stable/accurate). Ideally a dynamic time step should be used, but as long as the speed of sound is high enough, it should dominate in the example you are required to solve. You can therefore use the following fixed time step for your initial testing:

$$\Delta t = 0.1 \frac{h}{c_0}$$

Other required relationships

Another variable that you need to choose in the system is the relationship between the initial particle spacing, Δx and the characteristic smoothing length, h . You can use:

$$h = 1.3 \Delta x$$

Once you have the simulator working you can test the impact of changing this proportionality.

The initial mass of your particles also depends upon the initial particle spacing:

$$m_i = \Delta x^2 \rho_0$$

Smoothing the Density/Pressure field

In SPH the integration of the continuity equation will result in not only variations in the density (and thus the pressure) based on the macroscopic velocity gradients, but also on local variations in particle spacing and velocity. This will eventually result in a very “rough” density and pressure distribution.

To get around this problem the density should be smoothed every 10 to 20 time steps. The following smoothing should be used:

$$\rho_i = \frac{\sum_{j=1}^n w(\mathbf{r}_{ij}, h)}{\sum_{j=1}^n \frac{w(\mathbf{r}_{ij}, h)}{\rho_j}}$$

Note that the sums MUST include the $i = j$ combinations.

Finding Neighbours

Virtually all of the above equations require summing over neighbouring particles. If we were to naively search over all possible neighbours for each of the points required in the calculation, then the overall complexity of the algorithm would be $O(N^2)$ in the number of particles, which would be prohibitively expensive for all but the smallest problems.

We can get around this problem by exploiting the fact that the kernel has compact support and that we therefore only need to search for particles that are within $2h$ of the particle that we are considering. Considering that the particles will also be quite evenly spaced, the best way to do this search is using a linked cell method.

The way this works is that the entire domain is divided into a grid of $2h$ wide cells. A time step (or half time step in the predictor-corrector case) starts with each particle being assigned to their appropriate cell. This involves looping over all the particles and thus has a complexity of $O(N)$. The calculation for each particle involves looping over all the particles in its cell, together with all the particles in the neighbouring cells. Since the number of particles in these neighbouring cells does not depend upon the number of particles in the system, the calculations required to sum over all the neighbouring particles of a single particle is $O(1)$, with the calculations for all the particles thus being $O(N)$. The overall algorithm is thus $O(N)$.

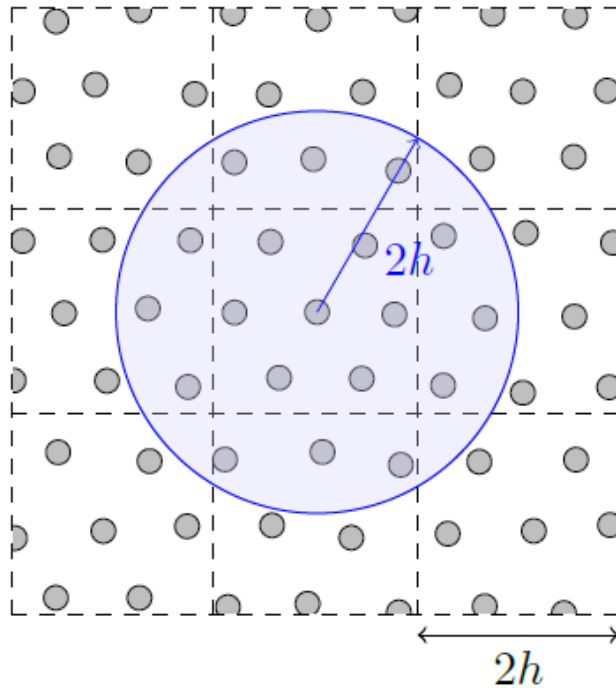


Figure 1: Searching using linked cell method

In order to aid the development of your program you have been given an implementation of the linked cell search method. This implementation is comparatively basic and does not include the ability to use different half-step locations. It will, though, show you the basics of how to implement this method and form a basis on which you can build and improve.

Boundaries and initial conditions in SPH

Being a Lagrangian method, domain boundaries can be problematic to implement in SPH. An easy way to do this is to put a layer of fixed particles $2h$ wide around the edge of the whole domain. These particles should interact with other particles in the same way as the fluid particles and should have a density (and thus also pressure) that is allowed to evolve in the same way as that of other particles. Their position should remain fixed, though, and their velocity should remain zero.

Despite the boundary particles possessing a pressure, fluid particles may occasionally leak out of the domain. The easiest thing to do is to delete them, but this results in a loss of fluid mass. You can push them back, but you need to watch out for what the appropriate density to give them is (if you don't they will clump as the local density will be higher than the average density assigned to the particle). The final alternative is to give the walls a repulsive force that gets high enough at short range that the particles never leak. This can be effective, but needs tuning and can result in instabilities if done incorrectly. Note that in order to ascertain if a particle has leaked you do need to know the location of the boundaries. Think about storing boundaries in their own data structures, as this will be especially useful for implementing arbitrary shaped boundaries.

The figure below gives the initial configuration that should be used for these simulations. In these simulations the following initial particle spacing is used:

$$\Delta x = 0.2 \text{ m}$$

Remember to make Δx a variable that can be changed in the simulation configuration so that you can easily carry out a convergence test. A warning when placing the points is that the code is likely to crash if there are 2 points that are either on top of one another or close to being on top of one another. I would therefore recommend writing a function which checks for concurrent or near concurrent points and run this function after the points have initially been place, but before the simulation runs. Make sure that you only remove one of any two near concurrent points! Also remember that if a fluid particle is near concurrent to a wall particle it is the fluid rather than the wall particle that must be removed.

The simulations will also be of water and so the appropriate physical properties should be used ($\rho_0 = 1000 \frac{kg}{m^3}$ and $\mu = 0.001 Pa s$).




Figure 2: Initial configuration for the simulation. Red are mobile fluid particles and blue are fixed boundary particles.

Code Structure

With your main simulation code you should separate it into a short front end in which the required fluid, geometry and simulator parameters are specified. This front end should call a script that contains all the numerical simulation code. This script should output data for all the particles at user prescribed simulation time intervals.

There should then be a separate program which loads these output files and performs the required visualisation and post-processing on them.

What is required of each team?

Tasks

- 1) Create an SPH simulator that can run based on the initial problem geometry and conditions described above. These simulations should run until $t = 30 \text{ s}$ (this may take a while under Python and so test the code over much shorter intervals). Start by implementing a simple forward Euler timestep method, for which you may need to use a very small timestep.
- 2) Output particle data to file in such a form that it can be used for subsequent post-processing. Remember that the time steps will be quite small and so only produce outputs at a user set simulation time output interval.
- 3) Create a post-processing code to plot the particle data stored in the output files produced by the SPH simulator. The output from this code should be an animation of the simulation results, saved in a sensible format.
- 4) Improve your timestepping method by implementing the predictor-corrector scheme described above. Note that you should save your functioning Forward Euler version as a separate file; you are not expected to have both timestepping methods implemented in the same code.
- 5) Measure the wave height and location of the crest as a function of time. What is the sloshing interval across the domain and how does this compare to the expected shallow water wave speed of $\sqrt{gh_0}$ where g is gravity and h_0 is the average initial water depth? Why might there be a discrepancy?
- 6) Carry out a convergence study in which the resolution (Δx) of the simulation is varied. You will need to decide upon a suitable metric for the error (for example, timescale for several wave transits). What is the order of convergence?
- 7) An “artificial pressure” contribution can improve the simulation results and aesthetics, especially when the fluid is in tension. J.J. Monaghan developed SPH and also introduced the use of an artificial pressure term. Look up his papers and implement this term.
- 8) Simulate different shaped domains, especially ones with non-vertical and horizontal walls. Try and simulate a shoaling wave (e.g. running up onto a beach).

Note that you are not expected to complete all of these tasks, though you should complete tasks 1-3 as a minimum. These tasks also carry the greatest weighting in terms of the marking scheme.

Recommendation for how to initially proceed

I would recommend that you quite rapidly decide upon a data structure for the particles. This can be based directly on the neighbour search code that you have been given for the Forward Euler timestepping. Once the data structure has been agreed, you can split the team into one group to implement the numerical solver and another group to handle data output and rendering. Code to write the data to file and the post-processing code can be developed without a functioning SPH simulator, using just the initial positions of the particles, which can be defined by the code given. It is important to get both aspects developed at the same time as you can't tell if the code is working correctly if you can't see the results.

In Task 4 you will need to modify this data structure to include the need to store two values of position, velocity and density for each particle in order to implement the predictor-corrector timestepping scheme, as well as modifying the neighbour search to use the appropriate particle positions. You do not need to modify the output routine or post-processing software for the predictor-corrector timestepping as only the full timestep state needs to be outputted.

Assessment

Your group project will be assessed in four ways:

Software (50 marks)

Software will be assessed based on functionality (25/50 marks) and sustainability (25/50 marks).

Functionality (25 marks): Your software will be assessed on its ability to perform the required tasks. Up to 15 marks will be awarded for successfully completing tasks 1-4; 10 marks being available for the more challenging later tasks. Marks will be deducted for (a) inaccuracy in the solution; (b) bugs or mistakes in implementation; (c) computational inefficiency (i.e., code efficiency is quite an important aspect of this project as it is a substantially more computationally demanding task than the previous one).

Sustainability (25 marks): As with all software projects, you should employ all the elements of best practice in software development that you have learned so far. A GitHub repository will be created for your project to host your software. The quality and sustainability of your software and its documentation will be assessed based on your final repository and how it evolves during the week. Please refer to the ACSE-4 handbook for more information about the assessment of software quality. As this project involves the development of a new simulator tool, you would not be expected to have a solution to test the code against for much of your development time. Test cases should therefore rather focus on the avoidance of diverging behaviour. Useful tests might include, for example, checking whether any particles exceed the speed of sound; checking whether particles have a density that is substantially (e.g., a factor of 1.5) different to the reference density; or checking whether particles are outside the simulation domain (the code should deal with particles that have leaked through the walls). Other important aspect of the sustainability of the project include sufficient and clear documentation of the code, clear descriptions of the code's usage, as well as examples and test cases.

Technical report (30 marks)

By the deadline (Friday 14th December, 12:00 noon), you must produce a technical report that presents the following elements, in the form of either a PDF or a Jupyter notebook (in your GitHub repo), accompanied by one or more videos:

1. A brief description of your solution algorithm. This must NOT include the description of the SPH method that you have been given in this document. It should rather focus on your specific implementation of the algorithm and the features that needed to be introduced to overcome some of the challenges encountered in producing a working SPH code. As part of this description you should include either a flowsheet or pseudocode which highlights the main simulation loops together with the simulation steps required within these loops. You are also advised to describe all steps taken to overcome numerical issues you encountered, such as particle leaking. (10 marks)

2. A demonstration execution of your software for prescribed conditions. This should ideally take the form of an animation or video. In the report you should include a few frames from the animation together with a brief description of the results. (5 marks)
3. A comparison between your software output in terms of average wave velocities as obtained from the location of the crest as a function of time and the expected shallow water wave speed from the approximation $\sqrt{gh_0}$. Discuss the discrepancies and their likely origin (2 marks)
4. Describe the results from the convergence study in which the resolution (Δx) of the simulation is varied. You should include a description of the error metric used and how it was obtained. Show a graph from which the order of convergence is obtained. We would not expect you to increase the resolution of the simulation significantly above that used in the base case due to the long simulation times when using Python. For similar reasons we would recommend using an error metric over a comparatively short simulation time (for instance over the time required for the wave to traverse the domain for the first time). (3 marks)
5. Demonstration of the impact of the use of an artificial pressure term on the results obtained (both quantitative and aesthetic) (3 marks).
6. Demonstrate results for simulations in a domain with boundaries that are at an angle other than vertical or horizontal (3 marks).
7. Layout and presentation of the Technical Report (4 marks).

Presentation (10 marks)

On Friday afternoon, you must present your software and analysis to the class. You will have 20 minutes, plus questions, for your presentation. Your presentation should cover all of the completed tasks. We recommend that your presentation is supported by slides made in Powerpoint or a similar tool. Your presentation should also cover a summary of how you managed the project and divided the tasks.

Teamwork (peer assessment; 10 marks)

After the presentations, you will complete a self-evaluation of your group's performance. This, together with observations made during the week by staff, will inform the teamwork component of your mark. Please refer to the ACSE-4 guidelines for more information about the assessment of teamwork.

Technical requirements

You should use the assigned GitHub repository exclusively for your project

Your software must be written in Python 3.6

You can import anything from the standard python 3.6 libraries as well as numpy, matplotlib and scipy, mpltools, mpl_toolkit, sympy (or your own implementations of algorithms available in scipy that you have learned in ACSE 3)

You are not allowed to import any other python packages (if in doubt, please query with the Module Coordinator)

You can assume that Users of your software will have pytest installed, so this does not need to be part of your repository

You should use Travis and the GitHub flow for any automated testing that you implement